

PLNmodels

A collection of Poisson lognormal models
for multivariate analysis of count data

J. Chiquet, M. Mariadassou, S. Robin, B. Batardière, J. Kwon

Paris-Saclay, AgroParisTech, INRAE

Last update 01 December, 2021

<https://pln-team.github.io/PLNmodels>

R/C++ Package PLNmodels

Last stable release on CRAN, development version available on GitHub.

```
install.packages("PLNmodels")  
remotes::install_github("PLN-team/PLNmodels@dev")
```

```
library(PLNmodels)  
packageVersion("PLNmodels")
```

```
## [1] '0.11.4'
```

Python module

A Python + PyTorch implementation is coming

Advertisement (more, sorry about that)

<https://computo.sfds.asso.fr>, a new journal promoting reproducible research

Help and documentation

The [PLNmodels website](#) contains the standard package documentation and a set of comprehensive vignettes for the top-level functions

Publications

Chiquet, J., M. Mariadassou, and S. Robin (2018). "Variational inference for probabilistic Poisson PCA". In: *The Annals of Applied Statistics* 12, pp. 2674-2698. URL: <http://dx.doi.org/10.1214/18-AOAS1177>.

Chiquet, J., M. Mariadassou, and S. Robin (2019). "Variational inference for sparse network reconstruction from count data". In: *Proceedings of the 19th International Conference on Machine Learning (ICML 2019)*.

Chiquet, J., M. Mariadassou, and S. Robin (2021). "The Poisson-Lognormal Model as a Versatile Framework for the Joint Analysis of Species Abundances". In: *Frontiers in Ecology and Evolution* 9. DOI: [10.3389/fevo.2021.588292](https://doi.org/10.3389/fevo.2021.588292).

Facon, B., A. Hafsi, M. C. de la Masselière, et al. (2021). "Joint species distributions reveal the combined effects of host plants, abiotic factors and species competition as drivers of species abundances in fruit flies". In: *Ecological Letters*. DOI: [10.1111/ele.13825](https://doi.org/10.1111/ele.13825).

PLNmodels: what is done¹

1. Motivations
2. A family of models for multivariate analysis
3. Efficient variational inference
4. Illustration

[1] and published

Routinely gathered in ecology/microbiology/genomics

Data tables

- **Abundances**: read counts of species/transcripts j in sample i
- **Covariates**: value of environmental variable k in sample i
- **Offsets**: sampling effort for species/transcripts j in sample i

Need frameworks to model *dependencies between counts*

- understand **environmental effects**
~> explanatory models (multivariate regression, classification)
- exhibit **patterns of diversity**
~> summarize the information (clustering, dimension reduction)
- understand **between-species interactions**
~> 'network' inference (variable/covariance selection)
- correct for technical and **confounding effects**
~> account for covariables and sampling effort

Models for multivariate count data

6 / 46

If we were in a Gaussian world...

The general linear model [MKB79] would be appropriate! For each sample $i = 1, \dots, n$,

$$\underbrace{\mathbf{Y}_i}_{\text{abundances}} = \underbrace{\mathbf{x}_i^\top \boldsymbol{\Theta}}_{\text{covariates}} + \underbrace{\mathbf{o}_i}_{\text{sampling effort}} + \boldsymbol{\epsilon}_i, \quad \boldsymbol{\epsilon}_i \sim \mathcal{N}(\mathbf{0}_p, \underbrace{\boldsymbol{\Sigma}}_{\text{between-species dependencies}})$$

null covariance \Leftrightarrow independence \rightsquigarrow uncorrelated species/transcripts do not interact

This model gives birth to Principal Component Analysis, Discriminant Analysis, Gaussian Graphical Models, Gaussian Mixture models and many others ...

With count data...

There is no generic model for multivariate counts

- Data transformation ($\log, \sqrt{\cdot}$): quick and dirty
- Non-Gaussian multivariate distributions [Ino+17]: do not scale to data dimension yet

The Poisson Lognormal model (PLN)

7 / 46

The PLN model [AH89] is a multivariate generalized linear model, where

- the counts \mathbf{Y}_i are the response variables
- the main effect is due to a linear combination of the covariates \mathbf{x}_i
- a vector of offsets \mathbf{o}_i can be specified for each sample.

$$\mathbf{Y}_i | \mathbf{Z}_i \sim \mathcal{P}(\exp \mathbf{Z}_i), \quad \mathbf{Z}_i \sim \mathcal{N}(\mathbf{o}_i + \mathbf{x}_i^\top \boldsymbol{\Theta}, \boldsymbol{\Sigma}),$$

The unknown parameters are

- $\boldsymbol{\Theta}$, the regression parameters
- $\boldsymbol{\Sigma}$, the variance-covariance matrix

Stacking all individuals together,

- \mathbf{Y} is the $n \times p$ matrix of counts
- \mathbf{X} is the $n \times d$ matrix of design
- \mathbf{O} is the $n \times p$ matrix of offsets

Properties: over-dispersion, arbitrary-signed covariances

- mean: $\mathbb{E}(Y_{ij}) = \exp(o_{ij} + \mathbf{x}_i^\top \boldsymbol{\Theta}_{.j} + \sigma_{jj}/2) > 0$
- variance: $\mathbb{V}(Y_{ij}) = \mathbb{E}(Y_{ij}) + \mathbb{E}(Y_{ij})^2 (e^{\sigma_{jj}} - 1) > \mathbb{E}(Y_{ij})$
- covariance: $\text{Cov}(Y_{ij}, Y_{ik}) = \mathbb{E}(Y_{ij})\mathbb{E}(Y_{ik}) (e^{\sigma_{jk}} - 1).$

Various tasks of multivariate analysis

- Dimension Reduction: rank constraint matrix Σ .

$$\mathbf{Z}_i \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma = \mathbf{B}\mathbf{B}^\top), \quad \mathbf{B} \in \mathcal{M}_{pk} \text{ with orthogonal columns.}$$

- Classification: maximize separation between groups with means

$$\mathbf{Z}_i \sim \mathcal{N}(\boldsymbol{\mu}_k \mathbf{1}_{\{i \in k\}}, \Sigma), \quad \text{for known memberships.}$$

- Clustering: mixture model in the latent space

$$\mathbf{Z}_i \mid i \in k \sim \mathcal{N}(\boldsymbol{\mu}_k, \Sigma_k), \quad \text{for unknown memberships.}$$

- Network inference: sparsity constraint on inverse covariance.

$$\mathbf{Z}_i \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma = \boldsymbol{\Omega}^{-1}), \quad \|\boldsymbol{\Omega}\|_1 < c.$$

- Variable selection: sparsity constraint on regression coefficients

$$\mathbf{Z}_i \sim \mathcal{N}(\mathbf{x}_i^\top \boldsymbol{\Theta}, \Sigma), \quad \|\boldsymbol{\Theta}\|_1 < c.$$

Inference: latent model but intractable EM

Estimate $\theta = (\Theta, \Sigma)$, predict the \mathbf{Z}_i , while the model marginal likelihood is

$$p_{\theta}(\mathbf{Y}_i) = \int_{\mathbb{R}_p} \prod_{j=1}^p p_{\theta}(Y_{ij} | Z_{ij}) p_{\theta}(\mathbf{Z}_i) d\mathbf{Z}_i$$

Maximum likelihood for incomplete data model: EM

With $\mathcal{H}(p) = -\mathbb{E}_p(\log(p))$ the entropy of p ,

$$\log p_{\theta}(\mathbf{Y}) = \mathbb{E}_{p_{\theta}(\mathbf{Z} | \mathbf{Y})}[\log p_{\theta}(\mathbf{Y}, \mathbf{Z})] + \mathcal{H}[p_{\theta}(\mathbf{Z} | \mathbf{Y})]$$

EM requires to evaluate (some moments of) $p_{\theta}(\mathbf{Z} | \mathbf{Y})$, but there is no close form!

Solutions

- [AH89] resort on numerical integration; [Kar05] Monte-Carlo integration
- Several heuristics, not always well motivated, found in the literature...
- Variational approach [WJ08]: use a proxy of $p_{\theta}(\mathbf{Z} | \mathbf{Y})$.

Principle

- Find a proxy of the conditional distribution $p(\mathbf{Z} \mid \mathbf{Y})$:

$$q(\mathbf{Z}) \approx p_{\theta}(\mathbf{Z} \mid \mathbf{Y}).$$

- Choose a convenient class of distribution \mathcal{Q} and minimize a divergence

$$q(\mathbf{Z})^* \arg \min_{q \in \mathcal{Q}} D(q(\mathbf{Z}), p(\mathbf{Z} \mid \mathbf{Y})).$$

Popular choice

The Küllback-Leibler divergence (error averaged wrt the approximated distribution)

$$KL(q(\mathbf{Z}), p(\mathbf{Z} \mid \mathbf{Y})) = \mathbb{E}_q \left[\log \frac{q(z)}{p(z)} \right] = \int_{\mathcal{Z}} q(z) \log \frac{q(z)}{p(z)} dz.$$

Class of distribution: diagonal multivariate Gaussian

$$\mathcal{Q} = \left\{ q : q(\mathbf{Z}) = \prod_i q_i(\mathbf{Z}_i), \quad q_i(\mathbf{Z}_i) = \mathcal{N}(\mathbf{Z}_i; \mathbf{m}_i, \text{diag}(\mathbf{s}_i \circ \mathbf{s}_i)), \mathbf{m}_i, \mathbf{s}_i \in \mathbb{R}_p \right\}$$

Maximize the ELBO (Evidence Lower BOund):

$$J(\theta, q) = \log p_\theta(\mathbf{Y}) - KL[q_\theta(\mathbf{Z}) || p_\theta(\mathbf{Z} | \mathbf{Y})] = \mathbb{E}_q[\log p_\theta(\mathbf{Y}, \mathbf{Z})] + \mathcal{H}[q(\mathbf{Z})]$$

Variational EM

- VE step: find the optimal q (here, $\{(\mathbf{m}_i, \mathbf{s}_i)\}_{i=1, \dots, n} = \{\mathbf{M}, \mathbf{S}\}$):

$$q^h = \arg \max J(\theta^h, q) = \arg \min_{q \in \mathcal{Q}} KL[q(\mathbf{Z}) || p_{\theta^h}(\mathbf{Z} | \mathbf{Y})]$$

- M step: update $\hat{\theta}^h$

$$\theta^h = \arg \max J(\theta, q^h) = \arg \max_{\theta} \mathbb{E}_{q^h}[\log p_\theta(\mathbf{Y}, \mathbf{Z})]$$

$$\text{Let } \mathbf{A} = \mathbb{E}_q[\exp(\mathbf{Z})] = \exp\left(\mathbf{O} + \mathbf{M} + \frac{1}{2}\mathbf{S}^2\right)$$

Variational bound

$$\begin{aligned} J(\mathbf{Y}) &= \mathbf{1}_n^\top ([\mathbf{Y} \circ (\mathbf{O} + \mathbf{M}) - \mathbf{A} + \log(\mathbf{S})]) \mathbf{1}_p + \frac{n}{2} \log |\mathbf{\Omega}| \\ &\quad - \frac{1}{2} \text{trace} \left(\mathbf{\Omega} \left[(\mathbf{M} - \mathbf{X}\mathbf{\Theta})^\top (\mathbf{M} - \mathbf{X}\mathbf{\Theta}) + \text{diag}(\mathbf{1}_n^\top \mathbf{S}^2) \right] \right) + \text{cst.} \end{aligned}$$

M-step (Analytical)

$$\hat{\mathbf{\Theta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X} \mathbf{M}, \quad \hat{\mathbf{\Sigma}} = \frac{1}{n} (\mathbf{M} - \mathbf{X} \hat{\mathbf{\Theta}})^\top (\mathbf{M} - \mathbf{X} \hat{\mathbf{\Theta}}) + \frac{1}{n} \text{diag}(\mathbf{1}^\top \mathbf{S}^2)$$

Variational E-step (optimization)

$$\frac{\partial J(q)}{\partial \mathbf{M}} = (\mathbf{Y} - \mathbf{A} - (\mathbf{M} - \mathbf{X}\mathbf{\Theta})\mathbf{\Omega}), \quad \frac{\partial J(q)}{\partial \mathbf{S}} = \frac{1}{\mathbf{S}} - \mathbf{S} \circ \mathbf{A} - \mathbf{S} \mathbf{D}_{\mathbf{\Omega}}.$$

Application to the optimization of PLN models

Property of PLN variational approximation

The ELBO $J(\theta, q)$ is bi-concave, i.e.

- concave wrt $q = (\mathbf{M}, \mathbf{S})$ for given θ
- concave wrt $\theta = (\boldsymbol{\Sigma}, \boldsymbol{\Theta})$ for given q but **not jointly concave** in general.

Optimization

Gradient ascent for the set of variational parameters (\mathbf{M}, \mathbf{S})

Medium scale problems

- **algorithm:** conservative convex separable approximations Svanberg [Sva02]
- **implementation:** `NLopt` nonlinear-optimization library Johnson [Joh11]
- **initialization:** LM after log-transformation applied independently on each variables + concatenation of the regression coefficients + Pearson residuals

↪ Comfortable up to a thousand of sites ($n \approx 1000$), hundreds of species ($p \approx 100s$)

Oaks powdery mildew data set overview 14 / 46

Jakuschkin, Fievet, Schwaller, Fort, Robin, and Vacher [Jak+16] Study effects of the pathogen *E.Aphiltoïdes* (mildew) wrt bacterial and microbial communities

Species Abundances

- Microbial communities sampled on the surface of $n = 116$ oak leaves
- Communities sequenced and cleaned resulting in $p = 114$ OTUs (66 bacteria, 48 fungi).

Covariates and offsets

Characterize the samples and the sampling, most important being

- `tree`: Tree status with respect to the pathogen (susceptible, intermediate or resistant)
- `distTOground`: Distance of the sampled leaf to the base of the ground
- `orientation`: Orientation of the branch (South-West SW or North-East NE)
- `readsTOTfun`: Total number of ITS1 reads for that leaf
- `readsTOTbac`: Total number of 16S reads for that leaf

Abundance table (I)

15 / 46

```
data(oaks)
oaks$Abundance %>% as_tibble() %>%
  dplyr::select(1:10) %>%
  rmarkdown::paged_table()
```

b_OTU_1045	b_OTU_109	b_OTU_1093	b_OTU_11	b_OTU_112
<int>	<int>	<int>	<int>	<int>
0	0	0	6	146
0	0	0	0	0
0	0	0	2	0
0	0	0	1	1
0	0	0	4	1
0	0	0	77	2
0	0	0	21	2
0	0	0	27	4
0	0	0	7	42
0	0	0	7	2


1-10 of 116 rows | 1-5 of 10 columns

Previous **1** 2 3 4 5 6 ... 12 Next

Abundance table (II)

16 / 46

```
log(1 + oaks$Abundance) %>%  
  corrrplot::corrrplot(is.corr = FALSE,  
    addgrid.col = NA, tl.cex = .5, cl.pos = "n")
```



PLN with offsets and covariates (1)

17 / 46

Offset: modeling sampling effort

The predefined offset uses the total sum of reads, accounting for technologies specific to fungi and bacteria:

```
M01_oaks ← PLN(Abundance ~ 1 + offset(log(Offset)) , oaks)
```

Covariates: tree and orientation effects ('ANOVA'-like)

The `tree` status is a natural candidate for explaining a part of the variance.

- We chose to describe the tree effect in the regression coefficient (mean)
- A possibly spurious effect regarding the interactions between species (covariance).

```
M11_oaks ← PLN(Abundance ~ 0 + tree + offset(log(Offset)), oaks)
```

What about adding more covariates in the model, e.g. the orientation?

```
M21_oaks ← PLN(Abundance ~ 0 + tree + orientation + offset(log(Offset)), oaks)
```

PLN with offsets and covariates (2)


18 / 46

There is a clear gain in introducing the tree covariate in the model:

```
rbind(M01 = M01_oaks$criteria,  
      M11 = M11_oaks$criteria, M21 = M21_oaks$criteria) %>%  
  knitr::kable(format = "html")
```

	nb_param	loglik	BIC	ICL
M01	6669	-32252.14	-48102.98	-52169.64
M11	6897	-31524.16	-47916.91	-51644.03
M21	7011	-31438.58	-48102.29	-51727.13

Looking at the coefficients Θ associated with `tree` bring additional insights:


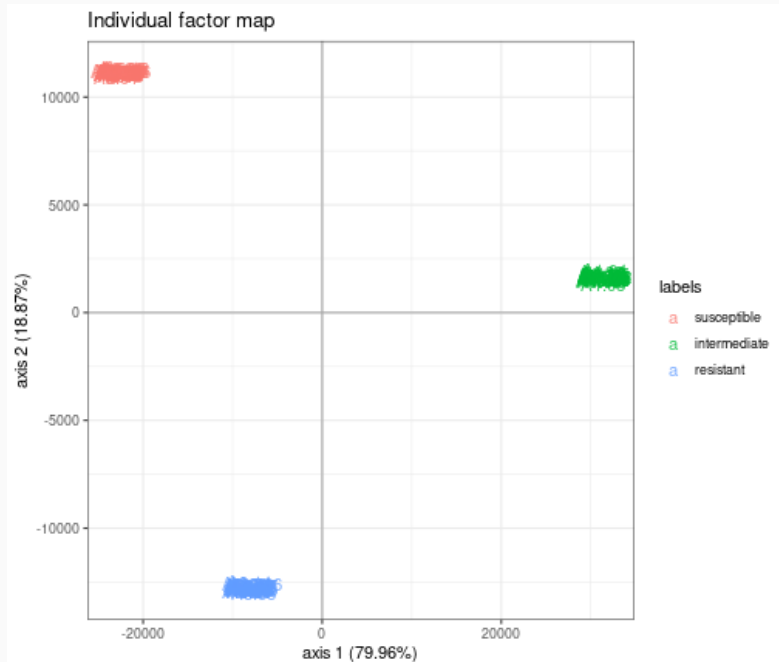


Discriminant Analysis


19 / 46

Use the `tree` variable for grouping (`grouping` is a factor of group to be considered)


```
myLDA_tree <-  
  PLNLDA(Abundance ~ 1 + offset(log(Offset)), grouping = oaks$tree, data = oaks)
```



```
PLNPCA(Abundance ~ 1 + offset(log(Offset)), data = oaks, ranks = 1:30)
PCA_offset_BIC ← getBestModel(PCA_offset, "BIC")
```



```
PCA_tree <-  
  PLNPCA(Abundance ~ 0 + tree + offset(log(Offset)), data = oaks, ranks = 1:30)
```



Clustering of the oaks samples

22 / 46

```
PLN_mixtures ←  
  PLNmixture(Abundance ~ 1 + offset(log(Offset)), data = oaks, clusters = 1:3)  
myPLN_mix ← getModel(PLN_mixtures, 3)
```

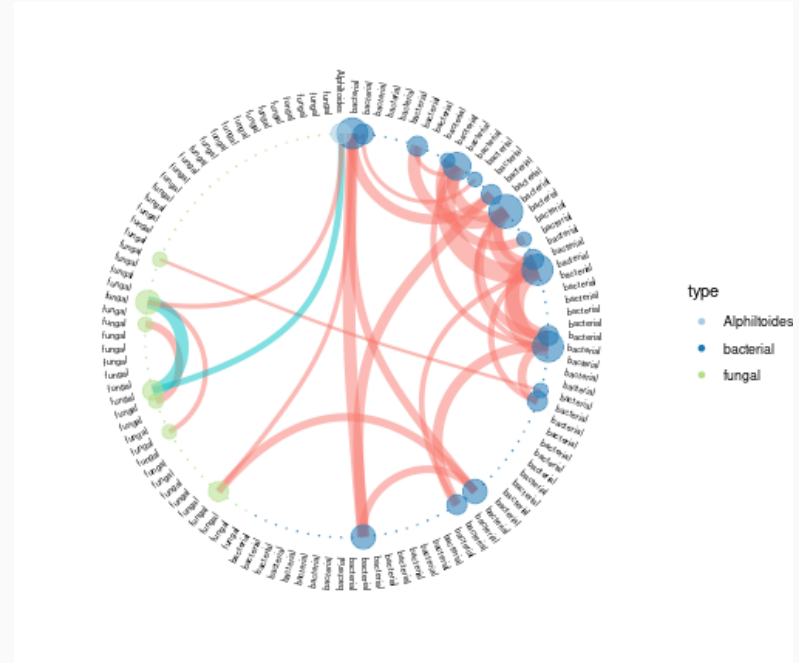

```
myPLN_mix$plot_clustering_pca()
```

```
myPLN_mix$plot_clustering_data()
```

Network inference

23 / 46

```
networks ← PLNnetwork(Abundance ~ 0 + tree + offset(log(Offset)), data = oaks)
```



Some more recent research

Limitations

- No guarantee for the variational estimators
- No test on the inference parameters
- Environmental Genomics, Single-Cell: more zeros, more rows, more columns

Ongoing work

- Hypothesis testing
- Zero-inflated version
- Large scale and/or exact optimization

We only consider the standard PLN model.

A first try: Wald test

Test $\mathcal{H}_0 : R\theta = r_0$ with the statistic

$$(R\hat{\theta} - r_0)^\top \left[nR\hat{\mathbb{V}}(\hat{\theta})R^\top \right]^{-1} (R\hat{\theta} - r_0) \sim \chi_k^2 \quad \text{where} \quad k = \text{rank}(R).$$

The Fisher Information matrix

$$I(\hat{\theta}) = -\mathbb{E}_\theta \left[\frac{\partial^2 \log \ell(\theta; x)}{\partial \theta^2} \right]$$

can be used as an approximation of $n\mathbb{V}(\hat{\theta})^{-1}$.

Application

Derive confidences interval for the inverse covariance $\mathbf{\Omega}$ and the regression parameters Θ .

Variational Fisher Information

The Fisher information matrix is given by

$$I(\theta) = \begin{pmatrix} \frac{1}{n}(\mathbf{I}_p \otimes \mathbf{X}^\top) \text{diag}(\text{vec}(\mathbf{A}))(\mathbf{I}_p \otimes \mathbf{X}) & \mathbf{0} \\ \mathbf{0} & \frac{1}{2}\mathbf{\Omega}^{-1} \otimes \mathbf{\Omega}^{-1} \end{pmatrix}$$

and can be inverted blockwise to estimate $\mathbb{V}(\hat{\theta})$.

Wald test and coverage


- $\hat{\mathbb{V}}(\Theta_{kj}) = [n(\mathbf{X}^\top \text{diag}(\text{vec}(\hat{A}_{\cdot j}))\mathbf{X})^{-1}]_{kk}$
- $\hat{\mathbb{V}}(\Omega_{kl}) = 2\Omega_{kk}\Omega_{ll}$

The confidence intervals at level α are given by

- $B_{kj} = \hat{B}_{kj} \pm \frac{q_{1-\alpha/2}}{\sqrt{n}} \sqrt{\hat{\mathbb{V}}(\Theta_{kj})}$
- $\Omega_{kl} = \hat{\Omega}_{kl} \pm \frac{q_{1-\alpha/2}}{\sqrt{n}} \sqrt{\hat{\mathbb{V}}(\Omega_{kl})}$.

Study Bias and coverage of the estimator


- number of samples $n \in \{50, 100, 500, 1000, 10000\}$
- number of species/genes $p \in \{20, 200\}$
- number of covariates $d \in \{2, 5, 10\}$
- sampling effort $TSS \in \{\text{low}, \text{medium}, \text{high}\}$ ($\approx 10^4$, 10^5 and 10^6)
- noise level $\sigma^2 \in \{0.2, 0.5, 1, 2\}$
- Σ as $\sigma_{jk} = \sigma^2 \rho^{|j-k|}$, with $\rho = 0.2$
- Θ with entries sampled from $\mathcal{N}(0, 1/d)$



The variational estimator are asymptotically unbiased

95% confident interval - Coverage

29 / 46



Variance underestimated, no trusted confidence intervals can be derived out-of-the box

Idea: M-estimation (with M. Mariadassou, S. Robin)

- Asymptotics for the VEM stationary point (not a log-likelihood stationary point).
- Use sandwich estimator for correction as in [WM15]

Motivations

- account for a large amount of zero, i.e. with single-cell data,
- try to separate "true" zeros from "technical"/dropouts

The Model

Use two latent vectors \mathbf{W}_i and \mathbf{Z}_i to model excess of zeroes and dependence structure

$$\begin{aligned}\mathbf{Z}_i &\sim \mathcal{N}(\mathbf{o}_i + \mathbf{x}_i^\top \boldsymbol{\Theta}, \boldsymbol{\Sigma}) \\ W_{ij} &\sim \mathcal{B}(\text{logit}^{-1}(\mathbf{x}_i^\top \boldsymbol{\Theta}_j^0)) \\ Y_{ij} | W_{ij}, Z_{ij} &\sim W_{ij}\delta_0 + (1 - W_{ij})\mathcal{P}(\exp\{Z_{ij}\}),\end{aligned}$$

The unknown parameters are

- $\boldsymbol{\Theta}$, the regression parameters (from the PLN component)
- $\boldsymbol{\Theta}^0$, the regression parameters (from the Bernoulli component)
- $\boldsymbol{\Sigma}$, the variance-covariance matrix

↪ ZI-PLN is a mixture of PLN and Bernoulli distribution with shared covariates.

Consider the standard ZIPLN model (*i.e.* not the ZIPLN-regression model) with 1 sample:

$$(W_j)_{j=1..p} \sim \mathcal{B}^{\otimes}(\pi) = \mathcal{B}(\pi_1) \otimes \dots \mathcal{B}(\pi_p)$$

$$(Z_j)_{j=1..p} \sim \mathcal{N}_p(\mu, \Sigma)$$

$$Y_j | W_j, Z_j \sim (1 - W_j)\mathcal{P}(e^{Z_j}) + W_j\delta_0$$

Proposition

The standard ZIPLN model defined above with parameter $\theta = (\pi, \mu, \Sigma)$ and parameter space $(0, 1)^p \times \mathbb{R}^p \times \mathbb{S}_p^{++}$ is identifiable.

Proof. We used the moments of \mathbf{Y} to prove identifiability and rely on the following results for Gaussian and Poisson distributions:

- If $U \sim \mathcal{N}(\mu, \sigma^2)$, then $\mathbb{E}[e^U] = \exp(\mu + \sigma^2/2)$
- If $U \sim \mathcal{P}(\lambda)$ then $\mathbb{E}[U] = \lambda \quad \mathbb{E}[U^2] = \lambda(1 + \lambda) \quad \mathbb{E}[U^3] = \lambda(1 + 3\lambda + \lambda^2)$

Each coordinate of θ can be expressed as a simple functions of the (first three) moments of p_θ and thus $p_\theta = p_{\theta'} \Rightarrow \theta = \theta'$.

Same routine...

Variational approximation

$$p(\mathbf{Z}_i, \mathbf{W}_i | \mathbf{Y}_i) \approx q_\psi(\mathbf{Z}_i, \mathbf{W}_i) \approx q_{\psi_1}(\mathbf{Z}_i) q_{\psi_2}(\mathbf{W}_i)$$

with

$$q_{\psi_1}(\mathbf{Z}_i) = \mathcal{N}(\mathbf{Z}_i; \mathbf{m}_i, \text{diag}(\mathbf{s}_i \circ \mathbf{s}_i)), \quad q_{\psi_2}(\mathbf{W}_i) = \bigotimes_{j=1}^p \mathcal{B}(W_{ij}, \pi_{ij})$$

Variational lower bound

Let $\theta = (\Theta, \Theta^0, \Sigma)$ and $\psi = (\mathbf{M}, \mathbf{S}, \Pi)$, then

$$\begin{aligned} J(\theta, \psi) &= \log p_\theta(\mathbf{Y}) - KL(p_\theta(\cdot | \mathbf{Y}) \| q_\psi(\cdot)) \\ &= \mathbb{E}_{q_\psi} \log p_\theta(\mathbf{Z}, \mathbf{W}, \mathbf{Y}) - \mathbb{E}_{q_\psi} \log q_\psi(\mathbf{Z}, \mathbf{W}) \\ &= \mathbb{E}_{q_\psi} \log p_\theta(\mathbf{Y} | \mathbf{Z}, \mathbf{W}) + \mathbb{E}_{q_{\psi_1}} \log p_\theta(\mathbf{Z}) + \mathbb{E}_{q_{\psi_2}} \log p_\theta(\mathbf{W}) \\ &\quad - \mathbb{E}_{q_{\psi_1}} \log q_{\psi_1}(\mathbf{Z}) - \mathbb{E}_{q_{\psi_2}} \log q_{\psi_2}(\mathbf{W}) \end{aligned}$$

Property: J is separately concave in θ , ψ_1 and ψ_2 .

A sparse criterion

Recall that $\theta = (\Theta, \Theta^0, \Omega = \Sigma^{-1})$. Sparsity allows to control the number of parameters:

$$\arg \min_{\theta, \psi} J(\theta, \psi) + \lambda_1 \|\Theta\|_1 + \lambda_2 \|\Omega\|_1 \left(+ \lambda_1 \|\Theta^0\|_1 \right)$$

Alternate optimization

- (Stochastic) Gradient-descent on $\Theta^0, \mathbf{M}, \mathbf{S}$
- Closed-form for posterior probabilities Π
- Inverse covariance Ω
 - if $\lambda_2 = 0$, $\hat{\Sigma} = n^{-1} \left[(\mathbf{M} - \mathbf{X}\Theta)^{\top} (\mathbf{M} - \mathbf{X}\Theta) + \bar{\mathbf{S}}^2 \right]$
 - if $\lambda_2 > 0$, ℓ_1 penalized MLE (\rightsquigarrow Graphical-Lasso with $\hat{\Sigma}$ as input)
- PLN regression coefficient Θ
 - if $\lambda_1 = 0$, $\hat{\Theta} = [\mathbf{X}^{\top} \mathbf{X}]^{-1} \mathbf{X}^{\top} \mathbf{M}$
 - if $\lambda_1 > 0$, vectorize and solve a ℓ_1 penalized least-squared problem

Initialize Θ^0 with logistic regression on $\delta_0(\mathbf{Y})$, Θ with Poisson regression

A quick example in genomics (1)

34 / 46

scRNA data set

The dataset `scRNA` contains the counts of the 500 most varying transcripts in the mixtures of 5 cell lines in human liver (obtained with standard 10x scRNAseq Chromium protocol).

We subsample 500 random cells and then keep the 200 most varying genes

```
library(ZIPLN)
data(scRNA)
scRNAsub      ← scRNA[sample.int(nrow(scRNA), 500), ]
scRNAsub$counts ← scRNAsub$counts[, 1:200]
scRNAsub$counts %>% as_tibble() %>% rmarkdown::paged_table()
```

KRT81 <int>	AKR1B10 <int>	LCN2 <int>	AKR1C2 <int>	ALDH1A1 <int>
211	123	1	60	36
3	2	2	1	0
1	4	285	10	0
2	1	1	0	2
7	2	190	2	3

A quick example in genomics (2)

35 / 46

Model fits

We adjust the standard PLN model and the ZI-PLN model with some sparsity on the precision matrix:

```
system.time(myPLN ←  
  PLN(counts ~ 1 + offset(log(total_counts)),  
    data = scRNAsub, control = list(trace = 0)))
```


```
##      user  system elapsed  
## 23.090    0.032   23.122
```

```
system.time(myZIPLN ←  
  ZIPLN(counts ~ 1 + offset(log(total_counts)), rho = .25,  
    data = scRNAsub, control = list(trace = 0)))
```

```
##      user  system elapsed  
## 50.279    0.052    7.572
```

A quick example in genomics (3)

36 / 46



ZI-PLN seems to be less variant for predicting small counts

A quick example in genomics (4)

37 / 46

```
prcomp(myZIPLN$latent) %>% factoextra::fviz_pca_ind(col.ind = scRNAsub$cell_line)
```



A quick example in genomics (5)

38 / 46

```
library(sbm)  
A ← myZIPLN$model_par$Omega ≠ 0; diag(A) ← 0  
mySBM ← estimateSimpleSBM(A, estimOptions=list(plot=FALSE))
```

A quick example in genomics (6)

39 / 46

```
library(igraph)  
G ← igraph::graph.adjacency(A, mode = "undirected")  
C ← igraph::cluster_fast_greedy(G)  
plot(C, G)
```


Large scale problems

Sophisticated Adaptive Stochastic Gradient Descent

- Rprop (1993) uses the gradient sign and update each variable independently:
- AdaGrad (2011) uses adaptive coordinate-wise step-sizes
- RMSProp (2012) adds momentum to the step-sizes
- Adam (2015) also adds momentum to the gradients



(a) SGD without momentum




(b) SGD with momentum


~> All available in **Pytorch** with auto-differentiation.

Optimizers comparison

41 / 46



$n = 1,000$, $p = 200$, $d = 2$. Rprop is much faster.



$n = 10,000, p = 2,000, d = 2$ (running time: 1 min 40s)

↪ Work up to $n = 100,000, p = 10,000s$

Exact likelihood maximization

43 / 46

With B. Batardière, J. Kwon

To compare and assess at least empirically the performance of the VE-M estimator

- Use importance sampling to estimate the likelihood:

$$p_{\theta}(Y_i) = \int \tilde{p}_{\theta}(Y_i|Z)p(Z)dZ = \int \tilde{p}_{\theta}(Z)dZ \approx \frac{1}{n_s} \sum_{k=1}^{n_s} \frac{\tilde{p}_{\theta}(V_k)}{g(V_k)}, \quad (V_k)_{1 \leq k \leq n_s} \stackrel{iid}{\sim} g$$

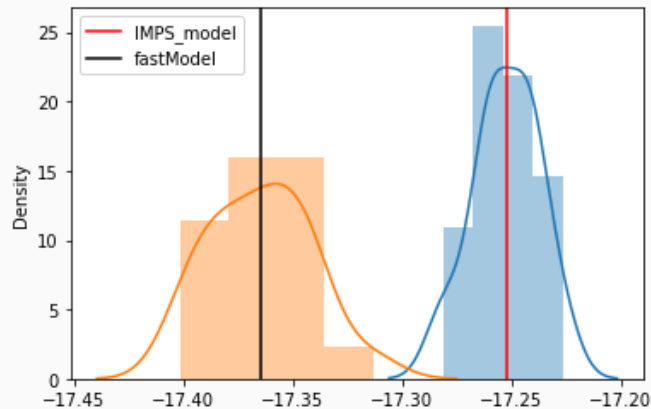
- Estimate the gradients of the log-likelihood by plug-in:

$$\nabla_{\theta} \log p_{\theta}(Y_i) \approx \nabla_{\theta} \log \left(\frac{1}{n_s} \sum_{k=1}^{n_s} \frac{\tilde{p}_{\theta}^{(u)}(V_k)}{g(V_k)} \right)$$


- Use SAGA (F. Bach), a smart incremental gradient descent approach

\rightsquigarrow Work up to $n = 1,000, p = 30, n = 1,000, q = 30, p \gg q$ for PCA variant

Comparing exact and variational estimators



Log-likelihood computed with importance sampling and variational estimators on 20 replicates



Same, on a single run (blue = IS, black = fastPLN, red = true log-likelihood)

The true log-likelihood used for comparison is computed numerically.

Summary

- PLN = generic model for multivariate count data analysis
- Flexible modeling of the covariance structure, allows for covariates
- Efficient V-EM algorithm

Extensions

- Other variants
 - covariance structures (spatial, time series, genetics...)
- Other models
 - Bernoulli/multinomial counterpart to PLN
 - functional data
 - multiple-data integration (e.g., Bernoulli + Poisson)

Aitchison, J. and C. Ho (1989). "The multivariate Poisson-log normal distribution". In: *Biometrika* 76.4, pp. 643-653.

Chiquet, J., M. Mariadassou, and S. Robin (2018). "Variational inference for probabilistic Poisson PCA". In: *The Annals of Applied Statistics* 12, pp. 2674-2698. URL: <http://dx.doi.org/10.1214/18-AOAS1177>.

Chiquet, J., M. Mariadassou, and S. Robin (2019). "Variational inference for sparse network reconstruction from count data". In: *Proceedings of the 19th International Conference on Machine Learning (ICML 2019)*.

Chiquet, J., M. Mariadassou, and S. Robin (2021). "The Poisson-Lognormal Model as a Versatile Framework for the Joint Analysis of Species Abundances". In: *Frontiers in Ecology and Evolution* 9. DOI: [10.3389/fevo.2021.588292](https://doi.org/10.3389/fevo.2021.588292).

Facon, B., A. Hafsi, M. C. de la Masselière, et al. (2021). "Joint species distributions reveal the combined effects of host plants, abiotic factors and species competition as drivers of species abundances in fruit flies". In: *Ecological Letters*. DOI: [10.1111/ele.13825](https://doi.org/10.1111/ele.13825).